Neuromorphic vs Photonic Computing for AI (2024–2025)

Executive overview

Recent 2024–2025 literature emphasizes neuromorphic computing as a promising route to sustain AI performance under energy and memory-bandwidth constraints, with digital CMOS platforms such as Intel Loihi and BrainScaleS-2 considered relatively mature and newer architectures like IBM NorthPole redefining memory–compute integration. In parallel, neuromorphic photonics has advanced rapidly, with Nature-family and other peer-reviewed works demonstrating photonic neuromorphic accelerators and lifelong-learning optical neural networks that can deliver order-of-magnitude energy-efficiency gains for specific workloads relative to conventional electronic neural networks. Overall, neuromorphic electronics currently offer greater algorithmic and tooling maturity for practical AI, while photonic neuromorphic and hybrid photonic–electronic architectures are emerging as ultra-high-throughput accelerators, particularly for dense linear operations.[1][2][3][4][5][6][7][8]

Recent neuromorphic roadmaps and surveys (2023–2025)

A major 2024 roadmap, "Roadmap to Neuromorphic Computing with Emerging Technologies," provides a comprehensive, materials-centric perspective on neuromorphic computing spanning CMOS, non-volatile memories, spintronics, photonics, and 2D materials. This roadmap classifies fully CMOS-based neuromorphic processors (TrueNorth, Intel Loihi, SpiNNaker, BrainScaleS) as the current large-scale hardware baseline and then surveys resistive memories (RRAM, PCM, FeRAM, MRAM), emerging memristive devices, and photonic components as enablers for future in-memory and brain-inspired computing systems. It highlights co-location of memory and compute, spiking/event-based processing, and mixed analog–digital designs as common themes, and explicitly flags photonic components plus phase-change and III–V materials as promising for neuromorphic photonic processors.[1][5]

The 2024 National Science Review article "Digital neuromorphic technology: current and future prospects" focuses on digital neuromorphic systems, emphasizing that digital approaches can leverage advanced technology nodes and provide repeatable behavior, while analog neuromorphic circuits have intrinsic energy-efficiency advantages but are harder to scale and integrate. This review identifies Intel's Loihi/Loihi 2 and SpiNNaker/SpiNNaker2 as representative digital platforms and concludes that digital neuromorphic hardware is "good to go," with adoption now limited more by algorithms, software tools, and programming abstractions than by hardware capabilities. It frames likely application domains as energy-efficient edge inference, event-based sensing, neurorobotics, and hybrid systems that combine spiking with conventional artificial neural networks.[^7]

On the photonic side, the 2023 "Neuromorphic Photonics Circuits: Contemporary Review" synthesizes progress in neuromorphic photonic integrated circuits, noting that optical implementations can perform complex computations with very high speed and "impressive energy efficiency" by exploiting parallelism, low loss, and high bandwidth. This review covers photonic neurons and synapses, reservoir computing, spiking and non-spiking optical neural networks, and discusses challenges such as integrating nonlinear elements, on-chip memory, and scalable training mechanisms. Nature's 2024 cross-journal collection "Neuromorphic hardware and computing" further aggregates recent research on neuromorphic architectures and devices, spanning both electronic and photonic realizations.[9][8] [^10]

Status of key platforms: Loihi, NorthPole, BrainScaleS-2

Architectural snapshot table

Platform	Type & implementation	Scale & performance (representative)	AI focus and maturity
Intel Loihi 2 / Hala Point	Digital spiking neuromorphic CMOS; event-driven SNN cores with on-chip learning and tightly integrated memory.[1][7]	Loihi 2 integrates up to ~1 million neurons per chip; Hala Point system combines 1,152 Loihi 2 chips for 1.15 billion neurons and 128 billion synapses in a 2.6 kW chassis, reaching up to ~15 TOPS/W and showing up to 100× lower energy and 50× higher speed than CPU/GPU baselines on certain real-time workloads.[11][12] [^13]	Research-grade but relatively mature; supports spiking and some conventional DNN workloads for edge and real-time AI (video, speech, wireless), with open-source Lava software stack and large external user community.[12][11] [^7]
IBM NorthPole	Digital in-memory "spatial computing" architecture; co-locates 224 MB of SRAM with 256 programmable cores, eliminating off-chip memory for inference.[14][6]	Fabricated in 12 nm with ~22 billion transistors; on ResNet-50 and YOLOv4, delivers ~25× higher energy efficiency than 12 nm GPUs and 14 nm CPUs and about 5× better frames-per-joule than NVIDIA H100, while maintaining competitive accuracy at 8/4/2-bit precision.[6][15]	Targeted at high-throughput neural inference (vision and, in newer work, 3-billion-parameter LLMs), currently used as a research prototype but demonstrating datacenter-level potential.[16][6]
BrainScaleS-2	Mixed-signal analog/digital neuromorphic system-on-chip; accelerated analog neurons and synapses with embedded digital processors and event routing.[17][18]	Each chip integrates 512 adaptive analog neurons and 131k plastic synapses; the system emulates neural dynamics about 1,000× faster than biological real time and is used in high-speed closed-loop robotics experiments.[19][17]	Developed mainly for computational neuroscience and spike-based machine learning; supports "hybrid plasticity" (analog measurements plus digital learning logic) and has been deployed via EBRAINS as a general research platform.[20][19]

Loihi 2 and Hala Point in more detail

The neuromorphic roadmap and Intel documentation describe Loihi as a fully digital spiking neuromorphic chip implementing programmable neuron models, on-chip learning rules, and event-driven communication with tightly coupled memory for synapses. Loihi 2, manufactured in Intel 4 process technology, scales this approach to roughly one million neurons per chip and underpins large systems such as Pohoiki Springs and the newer Hala Point neuromorphic supercomputer. Intel reports that Loihi-based systems can perform AI inference and optimization tasks using up to 100× less energy and up to 50× lower latency than conventional CPU/GPU architectures on suitable workloads, with Hala Point achieving efficiencies around 15 TOPS/W without input batching.[1][12] [11][13] [^7]

IBM NorthPole as an in-memory "neuromorphic" accelerator

IBM's NorthPole departs from classical von Neumann design by storing all model weights on-chip and intertwining memory and compute across 256 cores, each paired with local SRAM and connected via multiple on-chip networks. A 2023 Science article shows that, on ResNet-50 and YOLOv4 vision benchmarks, NorthPole is about 25× more energy-efficient than contemporary 12 nm GPUs and 14 nm CPUs and still about 5× more efficient than a 4 nm NVIDIA H100 GPU when measured in frames per joule. Recent IBM work extends this to a 3-billion-parameter LLM derived from Granite-8B, where NorthPole achieves sub-millisecond token latencies while maintaining a large energy-efficiency advantage over GPUs, highlighting its suitability for high-throughput inference rather than training.[14][16] [^6]

BrainScaleS-2 as an accelerated analog neuromorphic platform

The BrainScaleS-2 system implements physical analog neuron and synapse circuits on silicon that emulate adaptive-exponential neuron dynamics and plastic synapses in continuous time, combined with embedded digital processors for control and learning. Each BrainScaleS-2 core contains an analog crossbar of synapses and configurable neuron circuits, enabling both accelerated spiking dynamics (roughly 1,000× faster than biology) and support for non-spiking vector–matrix multiplications. Recent work demonstrates a high-speed robotics loop by tightly coupling BrainScaleS-2 to external sensors and actuators via a real-time event interface, leveraging the acceleration to realize microsecond-scale closed-loop control.[19][20] [^17]

Energy efficiency: neuromorphic electronics

The 2024 digital neuromorphic review concludes that analog neuromorphic circuits should, in principle, achieve higher energy efficiency than digital ones, but that digital neuromorphic chips benefit from advanced CMOS nodes and fabrication maturity, making both approaches relevant depending on application and development constraints. Across platforms such as Loihi 2 and NorthPole, the dominant energy advantages arise from event-driven operation, sparse data movement, and in-memory computing rather than from spiking alone.[1][11] [6][7]

Intel reports Loihi-based systems solving inference and optimization problems with up to 100× lower energy and up to 50× lower latency than CPU/GPU baselines on emerging edge workloads, and Hala Point demonstrates up to 15 TOPS/W while processing real-time data streams without batching. A 2024 study on sensor fusion for autonomous driving benchmarks Loihi 2 as over 100× more energy-efficient than a CPU and nearly 30× more energy-efficient than a GPU on representative localization tasks, while also running faster. These gains are consistent with earlier Loihi 2 and Loihi-based benchmarking work that finds orders-of-magnitude improvements for sparse, temporal, and event-based tasks compared with von Neumann architectures.[21][22] [23][11] [^13]

For NorthPole, the Science paper attributes its energy efficiency to fully on-chip model storage, localization of memory and compute, and extensive use of low-precision integer arithmetic. Relative to a 4 nm NVIDIA H100 GPU, NorthPole still delivers roughly a 5× improvement in frames per joule on image classification benchmarks, even though NorthPole is fabricated in an older 12 nm node, underscoring the impact of architectural choices on energy efficiency. These digital neuromorphic or neuromorphic-inspired accelerators thus demonstrate 10–100× improvements in energy per inference on suitable workloads versus conventional CPUs/GPUs, though the gains are highly workload-dependent and not yet universal across all AI tasks.[11][6] [^13]

BrainScaleS-2 and similar analog neuromorphic platforms can, in principle, reach even lower energy per synaptic operation by encoding computation directly in the physics of analog circuits and exploiting accelerated dynamics. While system-level energy numbers are less frequently benchmarked against GPUs, BrainScaleS-2 is explicitly positioned as a resource-efficient platform for large-scale neural model emulation, and its use in high-speed robotics suggests viability for low-latency, low-energy closed-loop tasks. The 2024 roadmap expects further orders-of-magnitude efficiency gains if emerging memristive and analog in-memory devices can be scaled and integrated reliably in neuromorphic architectures, though it notes substantial challenges around variability, endurance, and peripheral overheads.[19][20] [17][1]

Energy efficiency: neuromorphic photonics

The contemporary review of neuromorphic photonic circuits emphasizes that photonic implementations can execute complex computations with very high speed and "impressive energy efficiency," due to negligible resistive losses in waveguides, massive wavelength-division multiplexing, and sub-nanosecond device dynamics. Photonic neuromorphic systems are particularly well suited for parallel matrix–vector multiplications and reservoir computing, where linear optical operations can dominate computational cost.[9][8]

A 2024 Nature-family paper introduces L^{2}ONN, a reconfigurable photonic neuromorphic architecture for lifelong multi-task learning that exploits spatial sparsity and multi-wavelength parallelism to implement tens of tasks (vision, speech, medical diagnosis) in a single optical network. Experiments show that L^{2}ONN achieves more than an order-of-magnitude higher energy efficiency than representative electronic artificial neural networks and about 14× larger capacity than previous optical neural networks while maintaining competitive task performance. This demonstrates that carefully designed photonic neuromorphic architectures can significantly outperform electronic counterparts in energy efficiency for multi-task inference.[^3]

Another Nature article reports a photonic-integrated neuromorphic accelerator for convolutional neural networks that uses spectrum slicing and time multiplexing to perform convolution operations in a single optical pass. When combined with a simple electronic neural network back-end, this photonic accelerator reduces total power consumption by roughly 30% compared with the equivalent standalone digital neural network, while preserving accuracy, supporting its role as a low-power accelerator sub-system. A 2026 Nature paper goes further by integrating capacitive analog electronic memory co-located with photonic computing units, demonstrating over 26× power savings relative to conventional SRAM plus DAC architectures for on-chip training and inference on tasks such as MNIST digit recognition.[2][4]

These results collectively suggest that at the level of core compute kernels (e.g., convolutions or matrix–vector products), neuromorphic photonic hardware can offer between roughly 1.3× and >10× energy savings over state-of-the-art electronic neural accelerators, with the upper end realized in architectures that tightly co-integrate analog memory and optics. However, the neuromorphic photonics review and related works note that comparing full system energy remains challenging, because laser sources, optical I/O, control electronics, and optical–electrical conversion overheads are often excluded or treated optimistically in experimental setups. As a result, system-level energy advantages of photonic neuromorphic computing over advanced electronic neuromorphic or GPU hardware, while promising, are not yet as comprehensively quantified.[4][8] [2][3] [^9]

AI workload suitability: neuromorphic vs photonic

Digital neuromorphic reviews identify several classes of AI workloads where neuromorphic hardware is already competitive or superior to GPUs/CPUs: event-based sensing (e.g., dynamic vision sensors, neuromorphic cochleae), closed-loop control and robotics, low-power edge inference, and certain combinatorial optimization problems. Loihi systems, for example, have been used for real-time video, speech, and wireless signal processing at the edge, demonstrating significant energy and latency advantages when exploiting sparsity and event-driven computation. BrainScaleS-2 has been used for fast neurorobotics and spike-based machine learning, taking advantage of its accelerated dynamics for microsecond-scale feedback control.[1][22] [11][19] [^7]

IBM's NorthPole is optimized for high-throughput inference on dense deep neural networks, including ResNet-50, YOLOv4, and more recently a 3-billion-parameter LLM derived from Granite-8B, where it achieves state-of-the-art energy-latency trade-offs relative to GPUs. Because NorthPole stores the entire model on-chip and uses scheduled spatial computing, it is particularly suited to inference workloads where models fit within its SRAM budget; training remains on external digital hardware. In the medium term, the 2024 digital neuromorphic review suggests that hybrid systems combining spiking and conventional neural networks, possibly on platforms like SpiNNaker2 and Tianjic, will be important for broader AI adoption.[16][6] [^7]

Photonic neuromorphic architectures excel at workloads dominated by linear operations, such as large matrix–vector multiplications in convolutional and fully connected layers, and are especially promising when latency and throughput are critical. The L^2ONN photonic architecture explicitly targets multi-task lifelong learning, demonstrating that optical networks can incrementally learn tens of tasks without catastrophic forgetting while maintaining high efficiency. The photonic CNN accelerator and analog-memory-assisted photonic neuromorphic system show that photonics can provide substantial benefits as accelerators embedded within larger machine-learning pipelines, where electronic components still handle non-linearities, control, precision-critical operations, and long-term storage.[2][3] [4][9] [^8]

Current neuromorphic photonic platforms are typically demonstrated on modest-scale benchmarks (e.g., MNIST, small image datasets, reinforcement-learning control tasks) and rely on hybrid electronic–photonic training loops, which limits their immediate suitability for very large-scale AI models compared with digital neuromorphic chips like Hala Point or NorthPole. Nonetheless, their ability to process many channels or tasks in parallel with low latency and high bandwidth suggests strong potential as datacenter or edge accelerators, particularly for inference workloads that can be reformulated to exploit optical parallelism. [3][4] [24][8] [^2]

Hybrid photonic–neuromorphic architectures

Several lines of work explicitly pursue hybrid architectures that combine neuromorphic principles with photonic hardware and electronic memory or control. The 2024 neuromorphic roadmap notes "increasing interest in architectures that can exploit photonics components for computing" within neuromorphic systems, especially via silicon photonics co-integrated with phase-change optical memories. The Neuropuls project, for example, envisions neuromorphic accelerators based on integrated photonics augmented with phase-change and III–V materials on CMOS-compatible platforms, targeting high-efficiency neuromorphic computing.[1][25]

The photonic CNN accelerator mentioned earlier is explicitly framed as a photonic neuromorphic accelerator that serves as a sub-system within larger hybrid machine learning schemes, offloading convolutional operations to optics while leaving remaining layers to electronics. The 2026 Nature work on neuromorphic photonic computing with on-chip capacitive analog memory goes further by monolithically integrating analog electronic memory with photonic computing units to eliminate energy-costly data movement between electronics and photonics, achieving over 26× power savings compared with SRAM-DAC baselines. These architectures embody neuromorphic principles (co-located compute and memory, analog computation, parallelism) while using both electronic and photonic domains. [2][4]

An Optica 2026 paper describes a large-scale programmable photonic spiking neural network implemented on a two-chip system (MZI mesh plus laser array) and tested within an opto-electronic hybrid computing setup for reinforcement-learning control tasks such as CartPole and Pendulum. At the device-architecture level, patent literature and technical reports describe 3D electronic-photonic integrated neuromorphic circuits (3D EPIC) that combine photonic-memristive synapses, photonic dendrites, and electronic somas to realize high-density, high-connectivity neuromorphic systems. Together, these efforts illustrate a broad design space of hybrid photonic–neuromorphic architectures, ranging from photonic accelerators attached to digital neuromorphic or GPU systems to fully co-integrated photonic-electronic neuromorphic chips.[24][26]

Practical implications and open challenges

Across 2024–2025 surveys and roadmaps, a recurring conclusion is that neuromorphic electronics and neuromorphic photonics are complementary rather than competing paradigms for future AI hardware. Digital neuromorphic platforms like Loihi 2, Hala Point, NorthPole, SpiNNaker2, and BrainScaleS-2 are relatively mature and can already deliver 10–100× energy savings on specific edge, event-based, or inference workloads, but they still require advances in algorithms, programming models, and software tooling for widespread adoption. Neuromorphic photonic systems, while less mature in tooling and system-scale integration, demonstrate compelling per-operation energy and latency advantages and are particularly suitable as accelerators for dense linear operations and multi-task inference.[1][2] [11][3] [4][6] [7][8]

Major open challenges include scalable and robust on-chip learning for both spiking and photonic neuromorphic systems, managing device variability and analog noise (especially in memristive and photonic devices), integrating large on-chip or near-chip memories without erasing energy gains, and developing software frameworks that can target heterogeneous neuromorphic and photonic hardware. Roadmaps also emphasize the need for standardized benchmarks that fairly compare neuromorphic electronic, photonic, and conventional accelerators at the system level, including all supporting components such as lasers, drivers, and memory, to avoid over-optimistic energy claims. Addressing these challenges will be crucial to determining how neuromorphic electronic, photonic, and hybrid architectures ultimately share roles across edge devices, embedded systems, and large datacenters in the coming decade.[9][6] [7][8] [^1]

References

1. Roadmap to Neuromorphic Computing with Emerging Technologies - The roadmap is organized into several thematic sections, outlining current computing challenges, dis...
2. Neuromorphic photonic computing with an electro-optic analog ... - In neuromorphic photonic systems, device operations are typically governed by analog signals, necess...
3. Photonic neuromorphic architecture for tens-of-task lifelong learning - We innovate a reconfigurable lifelong-learning optical neural network (L2ONN), for highly-integrat...
4. Photonic neuromorphic accelerator for convolutional neural ... - Nature - We experimentally validate a photonic-integrated neuromorphic accelerator that uses a hardware-frien...
5. [PDF] Roadmap to neuromorphic computing with emerging technologies - REVIEW ARTICLE | OCTOBER 21 2024. Roadmap to neuromorphic computing with emerging technologies. Adna...
6. Neural inference at the frontier of energy, space, and time - Science - NorthPole has high energy efficiency because of less data movement, which is enabled by a lack of of...

7. Digital neuromorphic technology: current and future prospects - Digital neuromorphic technology: current and future prospects Open Access ... energy-efficiency, tho...
8. Neuromorphic Photonics Circuits: Contemporary Review - PubMed - Using optics and photonics principles, neuromorphic devices can execute intricate computations swift...
9. Neuromorphic Photonics Circuits: Contemporary Review - PMC - This paper reviews the recent developments in neuromorphic photonic integrated circuits, application...
10. Neuromorphic Hardware and Computing 2024 - Nature - In this cross-journal collection, we aim to bring together cutting-edge research of neuromorphic arc...
11. Intel Builds World's Largest Neuromorphic System to Enable More ... - Hala Point, the industry's first 1.15 billion neuron neuromorphic system, builds a path toward more ...
12. Neuromorphic Computing and Engineering with AI | Intel® - Intel Labs' neuromorphic research goes beyond today's deep-learning algorithms by co-designing optim...
13. Accelerating Sensor Fusion in Neuromorphic Computing: A Case Study on Loihi-2 - ...Motional, and Comma2k19. Our research demonstrated that Loihi-2, using spiking
neural networks, s...
14. A Look at NorthPole - IBM - Neuromorphic Chip - NorthPole is a neural inference architecture that blurs the boundary between compute and memory by e...
15. IBM Research's new NorthPole AI chip - A new chip architecture points to faster, more energy-efficient AI. A new chip prototype from IBM Re...
16. IBM's NorthPole achieves new speed and efficiency milestones - In new experimental results, the brain-inspired NorthPole research prototype chip has achieved lower...
17. The BrainScaleS-2 Accelerated Neuromorphic System With Hybrid ... - It combines a custom analog accelerator core supporting the accelerated physical emulation of bio-in...
18. BrainScaleS-2 — Heidelberg University - Open Neuromorphic - The BrainScaleS-2 is an accelerated spiking neuromorphic system-on-chip integrating 512 adaptive int...
19. Closing the loop: High-speed robotics with accelerated ... - Frontiers - The BrainScaleS-2 system is an established analog neuromorphic platform with versatile applications ...
20. The BrainScaleS-2 Accelerated Neuromorphic System With Hybrid ... - Here we describe the second generation of the BrainScaleS neuromorphic architecture, emphasizing app...
21. Efficient Video and Audio processing with Loihi 2 - Loihi 2 is an asynchronous, brain-inspired research processor that
generalizes several fundamental e...
22. Neuromorphic on-chip reservoir computing with spiking neural network architectures - ...architectures that excel in
different scenarios. The resulting networks demonstrate a range of be...
23. Implementing and Benchmarking the Locally Competitive Algorithm on the Loihi 2 Neuromorphic Processor - Neuromorphic processors have garnered considerable interest in recent years
for their potential in e...

24. Photonic chips advance real-time learning in spiking neural systems - The two-chip system includes a 16-channel photonic neuromorphic chip with 272 trainable parameters, ...
25. [PDF] Neuromorphic architectures based on augmented silicon photonics ... - In this work, we discuss our vision for neuromorphic accelerators based on integrated photonics with...
26. Energy efficiency improvements with neuromorphic materials - Explore the revolutionary potential of neuromorphic computing materials that mimic brain function fo...